

## ▼ Scenario 3: CROP RECOMMENDATION SYSTEM

A state Agricultural Development Officer is studying the relationship between soil conditions and crop yields. They have access to a dataset of soil nutrient levels and a shapefile of agricultural land parcels. They want to identify areas with optimal soil conditions for specific crops and recommend suitable crop types for those areas. Develop a model to solve the problem using Python and also explain your solution with complete documentation. Upload the solution code and documentation in the GitHub Public Repository. Share the code, visualization, GitHub Link and other stuff in the Google Classroom.

## ▼ About the data

The information utilized in this project is generated by enhancing and merging diverse publicly accessible datasets from India, encompassing factors such as weather and soil. The dataset is available for access [here](#). Unlike intricate factors that impact crop yield, this data is relatively straightforward, featuring a limited yet valuable set of features.

Included in the dataset are soil attributes like Nitrogen, Phosphorous, Potassium, and pH values. Additionally, it incorporates essential environmental factors such as humidity, temperature, and rainfall specific to each crop. Context Precision agriculture is in trend nowadays. It helps the farmers to get informed decision about the farming strategy. Here, I present you a dataset which would allow the users to build a predictive model to recommend the most suitable crops to grow in a particular farm based on various parameters.

Context This dataset was build by augmenting datasets of rainfall, climate and fertilizer data available for India.

Data fields

N - ratio of Nitrogen content in soil

P - ratio of Phosphorous content in soil

K - ratio of Potassium content in soil

temperature - temperature in degree Celsius

humidity - relative humidity in %

ph - ph value of the soil

rainfall - rainfall in mm

## ▼ Importing libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib as plt
from sklearn.model_selection import train_test_split
```


## ▼ Loading and Displaying Dataset

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
dataset=pd.read_csv("/content/drive/MyDrive/cropdataset/Crop_recommendation.csv")
```

```
dataset.head()
```

	N	P	K	temperature	humidity	ph	rainfall	label	
0	90	42	43	20.879744	82.002744	6.502985	202.935536	rice	
1	85	58	41	21.770462	80.319644	7.038096	226.655537	rice	
2	60	55	44	23.004459	82.320763	7.840207	263.964248	rice	
3	74	35	40	26.491096	80.158363	6.980401	242.864034	rice	
4	78	42	42	20.130175	81.604873	7.628473	262.717340	rice	

▼ this is a classification problem where the output will be in a categorical form

```
dataset.shape
```

(2200, 8)

```
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2200 entries, 0 to 2199
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   N            2200 non-null   int64
1   P            2200 non-null   int64
2   K            2200 non-null   int64
3   temperature  2200 non-null   float64
4   humidity     2200 non-null   float64
5   ph           2200 non-null   float64
6   rainfall     2200 non-null   float64
7   label        2200 non-null   object
dtypes: float64(4), int64(3), object(1)
memory usage: 137.6+ KB
```



```
dataset.isnull().sum()
```

```
N            0
P            0
K            0
temperature  0
humidity     0
ph           0
rainfall     0
label        0
dtype: int64
```

```
dataset.duplicated().sum()
```



0

```
dataset.describe()
```

	N	P	K	temperature	humidity	ph	rainfall	
count	2200.000000	2200.000000	2200.000000	2200.000000	2200.000000	2200.000000	2200.000000	
mean	50.551818	53.362727	48.149091	25.616244	71.481779	6.469480	103.463655	
std	36.917334	32.985883	50.647931	5.063749	22.263812	0.773938	54.958389	
min	0.000000	5.000000	5.000000	8.825675	14.258040	3.504752	20.211267	
25%	21.000000	28.000000	20.000000	22.769375	60.261953	5.971693	64.551686	
50%	37.000000	51.000000	32.000000	25.598693	80.473146	6.425045	94.867624	
75%	84.250000	68.000000	49.000000	28.561654	89.948771	6.923643	124.267508	
max	140.000000	145.000000	205.000000	43.675493	99.981876	9.935091	298.560117	


▼ Exploring Dataset

```
corr=dataset.drop(['label'],axis=1).corr()
corr
```

	N	P	K	temperature	humidity	ph	rainfall	
N	1.000000	-0.231460	-0.140512	0.026504	0.190688	0.096683	0.059020	
P	-0.231460	1.000000	0.736232	-0.127541	-0.118734	-0.138019	-0.063839	
K	-0.140512	0.736232	1.000000	-0.160387	0.190859	-0.169503	-0.053461	
temperature	0.026504	-0.127541	-0.160387	1.000000	0.205320	-0.017795	-0.030084	
humidity	0.190688	-0.118734	0.190859	0.205320	1.000000	-0.008483	0.094423	
ph	0.096683	-0.138019	-0.169503	-0.017795	-0.008483	1.000000	-0.109069	
rainfall	0.059020	-0.063839	-0.053461	-0.030084	0.094423	-0.109069	1.000000	

```
sns.heatmap(corr,annot=True,cbar=True,cmap='coolwarm')
```

<Axes: >




```
dataset['label'].value_counts()
```

```
rice      100
maize     100
jute      100
cotton    100
coconut   100
papaya    100
orange    100
apple     100
muskmelon 100
watermelon 100
grapes    100
mango     100
banana    100
pomegranate 100
lentil    100
blackgram 100
mungbean  100
mothbeans 100
pigeonpeas 100
kidneybeans 100
chickpea  100
coffee   100
Name: label, dtype: int64
```


```
dataset['label'].value_counts().plot(kind='pie', autopct="%.2f")
```

<Axes: ylabel='label'>



```
sns.displot(dataset['N'])
```

<seaborn.axisgrid.FacetGrid at 0x7da9fc61ceb0>



```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(16,12))  
plt.subplot(3,2,1)  
sns.distplot(dataset['N'])
```

```
plt.subplot(3,2,2)  
sns.distplot(dataset['P'])
```

```
plt.subplot(3,2,3)  
sns.distplot(dataset['K'])
```

```
plt.subplot(3,2,4)  
sns.distplot(dataset['temperature'])
```

```
plt.subplot(3,2,5)  
sns.distplot(dataset['ph'])
```

```
plt.subplot(3,2,6)  
sns.distplot(dataset['humidity'])
```

```
plt.show()
```

```
<ipython-input-163-ee3f2516cc75>:5: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(dataset['N'])
```

```
<ipython-input-163-ee3f2516cc75>:8: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(dataset['P'])
```

```
<ipython-input-163-ee3f2516cc75>:11: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(dataset['K'])
```

```
<ipython-input-163-ee3f2516cc75>:14: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(dataset['temperature'])
```

```
<ipython-input-163-ee3f2516cc75>:17: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(dataset['ph'])
```





```
<ipython-input-163-ee3f2516cc75>:20: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(dataset['humidity'])
```



```

import seaborn as sns
plt.figure(figsize=(16,10))
plt.subplot(3,2,1)
sns.boxplot(dataset['N'])

plt.subplot(3,2,2)
sns.boxplot(dataset['P'])

plt.subplot(3,2,3)
sns.boxplot(dataset['K'])







plt.subplot(3,2,4)
sns.boxplot(dataset['temperature'])

plt.subplot(3,2,5)
sns.boxplot(dataset['ph'])

plt.subplot(3,2,6)
sns.boxplot(dataset['humidity'])

plt.show()

```



#### ▼ Converting Categorical variables to a integer format

```
crop_dict = {
    'rice': 1,
    'maize': 2,
    'jute': 3,
    'cotton': 4,
    'coconut': 5,
    'papaya': 6,
    'orange': 7,
    'apple': 8,
    'muskmelon': 9,
    'watermelon': 10,
    'grapes': 11,
    'mango': 12,
    'banana': 13,
    'pomegranate': 14,
    'lentil': 15,
    'blackgram': 16,
    'mungbean': 17,
    'mothbeans': 18,
    'pigeonpeas': 19,
    'kidneybeans': 20,
    'chickpea': 21,
    'coffee': 22
}
dataset['crop_num']=dataset['label'].map(crop_dict)
```

```
dataset.head()
```

	N	P	K	temperature	humidity	ph	rainfall	label	crop_num	
0	90	42	43	20.879744	82.002744	6.502985	202.935536	rice	1	
1	85	58	41	21.770462	80.319644	7.038096	226.655537	rice	1	
2	60	55	44	23.004459	82.320763	7.840207	263.964248	rice	1	
3	74	35	40	26.491096	80.158363	6.980401	242.864034	rice	1	
4	78	42	42	20.130175	81.604873	7.628473	262.717340	rice	1	

#### ▼ Removing Label column from the dataset

```
dataset=dataset.drop('label',axis=1)
```

#### ▼ Training , Testing and Splitting Dataset

```
features=dataset.drop('crop_num',axis=1)
target=dataset['crop_num']
```

```
features.head()
```

	N	P	K	temperature	humidity	ph	rainfall	
0	90	42	43	20.879744	82.002744	6.502985	202.935536	
1	85	58	41	21.770462	80.319644	7.038096	226.655537	
2	60	55	44	23.004459	82.320763	7.840207	263.964248	
3	74	35	40	26.491096	80.158363	6.980401	242.864034	
4	78	42	42	20.130175	81.604873	7.628473	262.717340	

```
target.head()
```

```
0    1
1    1
2    1
3    1
4    1
Name: crop_num, dtype: int64
```

```
# Splitting into train and test data
```

```
from sklearn.model_selection import train_test_split
Xtrain, Xtest, Ytrain, Ytest = train_test_split(features,target,test_size = 0.2,random_state =2)
```

Xtrain.shape

(1760, 7)

Xtest.shape

(440, 7)

Xtrain

	N	P	K	temperature	humidity	ph	rainfall
1936	113	38	25	22.000851	79.472710	7.388266	90.422242
610	28	35	22	29.530376	86.733460	7.156563	59.872321
372	11	61	21	18.623288	23.024103	5.532101	135.337803
1559	29	139	205	23.641424	93.744615	6.155939	116.691218
1500	24	128	196	22.750888	90.694892	5.521467	110.431786
...	...	...	...	...	...	...	...
1071	105	88	54	25.787498	84.511942	6.020445	114.200546
433	27	71	23	23.453790	46.487148	7.109598	150.871220
674	23	39	22	29.256493	81.979522	6.864839	42.024833
1099	117	81	53	29.507046	78.205856	5.507642	98.125658
1608	39	24	14	30.554726	90.903438	7.189260	106.071198

1760 rows × 7 columns

```
# Initializing empty lists to append all model's name and corresponding name
acc = []
model = []
```

Model Building

Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report
from sklearn import metrics
from sklearn import tree
DecisionTree = DecisionTreeClassifier(criterion="entropy",random_state=2,max_depth=5)

DecisionTree.fit(Xtrain,Ytrain)

predicted_values = DecisionTree.predict(Xtest)
x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Decision Tree')
print("DecisionTrees's Accuracy is: ", x*100)

print(classification_report(Ytest,predicted_values))
```

DecisionTrees's Accuracy is: 90.0					
	precision	recall	f1-score	support	
1	1.00	0.62	0.77	16	
2	1.00	1.00	1.00	21	
3	0.74	0.93	0.83	28	
4	1.00	1.00	1.00	20	
5	0.91	1.00	0.95	21	
6	1.00	0.84	0.91	19	
7	1.00	1.00	1.00	29	
8	1.00	1.00	1.00	13	
9	1.00	1.00	1.00	23	
10	1.00	1.00	1.00	15	
11	1.00	1.00	1.00	18	
12	1.00	1.00	1.00	26	
13	1.00	1.00	1.00	17	
14	1.00	1.00	1.00	17	
15	0.68	1.00	0.81	23	
16	0.59	1.00	0.74	16	
17	1.00	1.00	1.00	24	
18	0.00	0.00	0.00	19	



19	0.62	1.00	0.77	18
20	0.00	0.00	0.00	14
21	1.00	1.00	1.00	21
22	1.00	1.00	1.00	22
accuracy			0.90	440
macro avg	0.84	0.88	0.85	440
weighted avg	0.86	0.90	0.87	440

```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are
_warn_prf(average, modifier, msg_start, len(result))

```

```
from sklearn.model_selection import cross_val_score
```

```

# Cross validation score (Decision Tree)
score = cross_val_score(DecisionTree, X, target, cv=5)

```

```
score
```

```
array([0.93636364, 0.90909091, 0.91818182, 0.85909091, 0.93636364])
```

## ▼ Saving trained Decision Tree model

```

import pickle
# Dump the trained Naive Bayes classifier with Pickle
DT_pkl_filename = 'DecisionTree.pkl'
# Open the file to save as pkl file
DT_Model_pkl = open(DT_pkl_filename, 'wb')
pickle.dump(DecisionTree, DT_Model_pkl)
# Close the pickle instances
DT_Model_pkl.close()

```

## ▼ Guassian Naive Bayes

```
from sklearn.naive_bayes import GaussianNB
```

```
NaiveBayes = GaussianNB()
```

```
NaiveBayes.fit(Xtrain, Ytrain)
```

```

predicted_values = NaiveBayes.predict(Xtest)
x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Naive Bayes')
print("Naive Bayes's Accuracy is: ", x)

```

```
print(classification_report(Ytest, predicted_values))
```

```

Naive Bayes's Accuracy is: 0.990909090909091
      precision    recall  f1-score   support

     1         1.00      0.75      0.86         16
     2         1.00      1.00      1.00         21
     3         0.88      1.00      0.93         28
     4         1.00      1.00      1.00         20
     5         1.00      1.00      1.00         21
     6         1.00      1.00      1.00         19
     7         1.00      1.00      1.00         29
     8         1.00      1.00      1.00         13
     9         1.00      1.00      1.00         23
    10         1.00      1.00      1.00         15
    11         1.00      1.00      1.00         18
    12         1.00      1.00      1.00         26
    13         1.00      1.00      1.00         17
    14         1.00      1.00      1.00         17
    15         1.00      1.00      1.00         23
    16         1.00      1.00      1.00         16
    17         1.00      1.00      1.00         24
    18         1.00      1.00      1.00         19
    19         1.00      1.00      1.00         18
    20         1.00      1.00      1.00         14
    21         1.00      1.00      1.00         21
    22         1.00      1.00      1.00         22

```

accuracy			0.99	440
macro avg	0.99	0.99	0.99	440
weighted avg	0.99	0.99	0.99	440

```
# Cross validation score (NaiveBayes)
score = cross_val_score(NaiveBayes, features, target, cv=5)
score

array([0.99772727, 0.99545455, 0.99545455, 0.99545455, 0.99090909])
```

## ▼ Saving trained Guassian Naive Bayes model

```
import pickle
# Dump the trained Naive Bayes classifier with Pickle
NB_pkl_filename = 'NBClassifier.pkl'
# Open the file to save as pkl file
NB_Model_pkl = open(NB_pkl_filename, 'wb')
pickle.dump(NaiveBayes, NB_Model_pkl)
# Close the pickle instances
NB_Model_pkl.close()
```

## ▼ Support Vector Machine (SVM)

```
from sklearn.svm import SVC

SVM = SVC(gamma='auto')

SVM.fit(Xtrain, Ytrain)

predicted_values = SVM.predict(Xtest)

x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('SVM')
print("SVM's Accuracy is: ", x)

print(classification_report(Ytest, predicted_values))

SVM's Accuracy is: 0.10681818181818181
      precision    recall  f1-score   support

     1         0.50         0.06         0.11         16
     2         0.00         0.00         0.00         21
     3         1.00         0.07         0.13         28
     4         1.00         0.05         0.10         20
     5         1.00         0.05         0.09         21
     6         1.00         0.05         0.10         19
     7         1.00         0.03         0.07         29
     8         1.00         0.23         0.38         13
     9         1.00         0.30         0.47         23
    10         1.00         0.13         0.24         15
    11         1.00         0.06         0.11         18
    12         0.00         0.00         0.00         26
    13         1.00         0.24         0.38         17
    14         1.00         0.12         0.21         17
    15         0.00         0.00         0.00         23
    16         1.00         0.19         0.32         16
    17         1.00         0.12         0.22         24
    18         0.00         0.00         0.00         19
    19         0.00         0.00         0.00         18
    20         0.03         1.00         0.07         14
    21         1.00         0.05         0.09         21
    22         0.00         0.00         0.00         22

      accuracy
macro avg      0.66      0.13      0.14      440
weighted avg    0.66      0.11      0.13      440
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are
_warn_prf(average, modifier, msg_start, len(result))
```

```
# Cross validation score (SVM)
score = cross_val_score(SVM, features, target, cv=5)
score

array([0.27727273, 0.28863636, 0.29090909, 0.275      , 0.26818182])
```

## ▼ Logistic Regression

```
from sklearn.linear_model import LogisticRegression

LogReg = LogisticRegression(random_state=2)

LogReg.fit(Xtrain, Ytrain)

predicted_values = LogReg.predict(Xtest)

x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('Logistic Regression')
print("Logistic Regression's Accuracy is: ", x)

print(classification_report(Ytest, predicted_values))
```

Logistic Regression's Accuracy is: 0.9522727272727273

	precision	recall	f1-score	support
1	0.85	0.69	0.76	16
2	0.90	0.86	0.88	21
3	0.84	0.93	0.88	28
4	0.86	0.90	0.88	20
5	1.00	1.00	1.00	21
6	1.00	0.95	0.97	19
7	1.00	1.00	1.00	29
8	1.00	1.00	1.00	13
9	1.00	1.00	1.00	23
10	1.00	1.00	1.00	15
11	1.00	1.00	1.00	18
12	0.96	1.00	0.98	26
13	1.00	1.00	1.00	17
14	1.00	1.00	1.00	17
15	0.88	1.00	0.94	23
16	0.86	0.75	0.80	16
17	1.00	0.96	0.98	24
18	0.84	0.84	0.84	19
19	1.00	1.00	1.00	18
20	1.00	1.00	1.00	14
21	1.00	1.00	1.00	21
22	1.00	1.00	1.00	22
accuracy			0.95	440
macro avg	0.95	0.95	0.95	440
weighted avg	0.95	0.95	0.95	440

/usr/local/lib/python3.10/dist-packages/sklearn/linear\_model/\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

```
# Cross validation score (Logistic Regression)
score = cross_val_score(LogReg, features, target, cv=5)
score
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear\_model/\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

/usr/local/lib/python3.10/dist-packages/sklearn/linear\_model/\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
array([0.95      , 0.96590909, 0.94772727, 0.96818182, 0.94318182])
```

## ▼ Saving trained Logistic Regression model

```
import pickle
# Dump the trained Naive Bayes classifier with Pickle
LR_pkl_filename = 'LogisticRegression.pkl'
# Open the file to save as pkl file
LR_Model_pkl = open(DT_pkl_filename, 'wb')
pickle.dump(LogReg, LR_Model_pkl)
# Close the pickle instances
LR_Model_pkl.close()
```

## ▼ Random Forest

```
from sklearn.ensemble import RandomForestClassifier

RF = RandomForestClassifier(n_estimators=20, random_state=0)
RF.fit(Xtrain,Ytrain)

predicted_values = RF.predict(Xtest)

x = metrics.accuracy_score(Ytest, predicted_values)
acc.append(x)
model.append('RF')
print("RF's Accuracy is: ", x)

print(classification_report(Ytest,predicted_values))
```

```
RF's Accuracy is: 0.990909090909091
precision    recall  f1-score   support

   1         1.00      0.81      0.90         16
   2         1.00      1.00      1.00         21
   3         0.90      1.00      0.95         28
   4         1.00      1.00      1.00         20
   5         1.00      1.00      1.00         21
   6         1.00      1.00      1.00         19
   7         1.00      1.00      1.00         29
   8         1.00      1.00      1.00         13
   9         1.00      1.00      1.00         23
  10         1.00      1.00      1.00         15
  11         1.00      1.00      1.00         18
  12         1.00      1.00      1.00         26
  13         1.00      1.00      1.00         17
  14         1.00      1.00      1.00         17
  15         1.00      1.00      1.00         23
  16         0.94      1.00      0.97         16
  17         1.00      1.00      1.00         24
  18         1.00      0.95      0.97         19
  19         1.00      1.00      1.00         18
  20         1.00      1.00      1.00         14
  21         1.00      1.00      1.00         21
  22         1.00      1.00      1.00         22

accuracy                   0.99         440
```

macro avg	0.99	0.99	0.99	440
weighted avg	0.99	0.99	0.99	440

```
# Cross validation score (Random Forest)
score = cross_val_score(RF, features, target, cv=5)
score

array([0.99545455, 0.99545455, 0.99545455, 0.99318182, 0.98863636])
```

### ▼ Saving trained Random Forest model

```
import pickle
# Dump the trained Naive Bayes classifier with Pickle
RF_pkl_filename = 'RandomForest.pkl'
# Open the file to save as pkl file
RF_Model_pkl = open(RF_pkl_filename, 'wb')
pickle.dump(RF, RF_Model_pkl)
# Close the pickle instances
RF_Model_pkl.close()
```

### ▼ Accuracy Comparison

```
accuracy_models = dict(zip(model, acc))
for k, v in accuracy_models.items():
    print (k, '=', v)

Decision Tree = 0.9
Naive Bayes = 0.990909090909091
SVM = 0.10681818181818181
Logistic Regression = 0.9522727272727273
RF = 0.990909090909091
```

### ▼ From the Above Accuracies, We have adopt Random Forest as it has Maximum Accuracy

```
classifier=RandomForestClassifier()
classifier.fit(Xtrain,Ytrain)
ypred=classifier.predict(Xtest)
```

```
accuracy_score(Ytest,ypred)*100

99.31818181818181
```

### ▼ Sample Predictions

```
def recommendation(N,P,k,temperature,humidity,ph,rainfal):
    features = np.array([[N,P,k,temperature,humidity,ph,rainfal]])
    prediction = DecisionTree.predict(features).reshape(1,-1)

    return prediction[0]
```